

An Introduction to PETSc: The Portable, Extensible Toolkit for Scientific Computation

A User's Overview

Jeff Borggaard

Virginia Tech

17 January 2006

Disclaimer

PETSc was developed at Argonne National Laboratory by Satish Balay, Kris Buschelman, Victor Eijkhout, William Gropp, Dinesh Kaushik, Matthew Knepley, Lois Curfman McInnes, Barry Smith and Hong Zhang.

tt <http://www.mcs.anl.gov/petsc>

Along with Traian Iliescu and Alexey Miroshnikov, I have *used* PETSc to build a parallel 3D finite element flow solver for evaluating LES models of realistic flows.

What is PETSc?

PETSc, the Portable, Extensible Toolkit for Scientific Computation:

- Data structures
- Functions

Built upon blas, linpack and MPI.

Built specifically to shorten the development time of parallel scientific application software. Especially on distributed memory architectures requiring message passing.

Why PETSc?

Why PETSc?

Functional

- Data structures
 - parallel vectors (sequential, MPI)
 - sparse matrices (compressed sparse row, etc.)
 - distributed arrays
 - index sets
- Functions
 - Krylov subspace methods (GMRES, CG, CGS, etc.)
 - preconditioners
 - nonlinear solvers (line search, trust-region)
 - time steppers (forward/backward Euler, pseudo-time stepping)
- Profiling
 - `-log_summary`

Portable

- Freely available (www.mcs.anl.gov/petsc)
- Very responsive developers (petsc-maint@mcs.anl.gov)
- Available for C, C++, FORTRAN 77/90
- Ported to a wide variety of platforms

Distributed Memory

- HP
- IBM
- SGI
- Sun
- Mac OS X
- PCs: Linux and Wintel

Shared Memory

- Cray T3E
- HP 9000
- IBM SP
- SGI Origin
- Sun Enterprise

Extensible

Has been coupled numerous packages, including

- ADIC: Automatic Differentiation in C
- Chaco, Jostle, MeTiS: Graph partitioning packages
- Mathematica
- MATLAB
- Trilinos: Multilevel preconditioning package

But ...

It has a (reasonbly) steep learning curve

How to use PETSc

Set up environment variables. On `phoenix.scs.fsu.edu`:

- `setenv PETSC_DIR /usr/local/petsc`
- `setenv PETSC_ARCH linux-gnu`

Add the following line to your makefile:

- `include ${PETSC_DIR}/bmake/common/base`

then link with eg. `${PETSC_SYS_LIB}` or `${PETSC_KSP_LIB}`

Run just as any MPI job

- `mpirun -np 4 petsc-exec (options)`

inside of a condor script.

Hello World! from MPI and PETSc

```
// MPI hello world
#include <stdio.h>
#include "mpi.h"

int main( int argc,
          char *argv[] ) {
    MPI_Init( &argc,
              &argv);
    printf( "Hello World\n");
    MPI_Finalize();
    return 0;
}
```

```
// PETSc hello world
#include "petsc.h"

int main( int argc,
          char *argv[] ) {
    PetscInitialize( &argc,
                    &argv,
                    PETSC_NULL,
                    PETSC_NULL);
    PetscPrintf( PETSC_COMM_WORLD,
                 "Hello World\n");
    PetscFinalize();
    return 0;
}
```

Hello World! from MPI and PETSc (2)

```
node001% mpirun -np 4 hello_mpi  
Hello World  
Hello World  
Hello World  
Hello World  
node001%
```

```
node001% mpirun -np 4 hello_petsc  
Hello World  
node001%
```

Hello World! from MPI and PETSc (3)

```
// The MPI hello world program (modified to print once)
#include <stdio.h>
#include "mpi.h"

int main( int argc, char *argv[] ) {
    int rank;

    MPI_Init( &argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        printf( "Hello World\n");
    }
    MPI_Finalize();
    return 0;
}
```

First Example

We are going to limit our discussion to four components from PETSc,

- Vectors (Vec*)
- Matrices (Mat*)
- Krylov Subspace Methods (KSP*)
- Preconditioners (PC*)

and use them to solve a linear PDE.

These are the essential building blocks of PETSc.

PETSc Vectors

There are essentially two types of PETSc Vectors:

- Sequential (local copy on each processor)

```
VecCreateSeq(PETSC_COMM_SELF, int m, Vec *x);
```

and

- Parallel (large vector stored over all processors)

```
VecCreateMPI(PETSC_COMM_WORLD, int m, int M, Vec *x);
```

a third, `VecCreateShared()`, could be used on a shared memory machine.

Either `m` or `M` above could be set to `PETSC_DECIDE` to have PETSc determine it.

PETSc Vectors (2)

Vectors can be filled using a number of functions. These include

- Setting the entire vector to one value (often 0)

```
VecSet(Vec x, PetscScalar value);
```

or to set individual components

- either *insert* values

```
VecSetValues(Vec x, int n, int *indices, PetscScalar  
*values, INSERT_VALUES);
```

- or *add* values

```
VecSetValues(Vec x, int n, int *indices, PetscScalar  
*values, ADD_VALUES);
```

Once all components have been set with `VecSetValues()`, the following functions must be called

- `VecAssemblyBegin(Vec x);`
- `VecAssemblyEnd(Vec x);`

PETSc Vectors (3)

A partial list of functions on PETSc Vectors follows:

<code>VecAXPY(Vec y, PetscScalar a, Vec x);</code>	$y = y + a*x$
<code>VecScale(Vec x, PetscScalar a);</code>	$x = a*x$
<code>VecDot(Vec x, Vec y, PetscScalar *r);</code>	$r = \bar{x}'*y$
<code>VecNorm(Vec x, NormType type, double *r);</code>	$r = \ x\ _{type}$
<code>VecSum(Vec x, PetscScalar *r);</code>	$r = \sum x_i$
<code>VecCopy(Vec x, Vec y);</code>	$y = x$
<code>VecPointwiseMult(Vec w, Vec x, Vec y);</code>	$w_i = x_i * y_i$
<code>VecMax(Vec x, int *index, double *r);</code>	$r = \max x_i$
<code>VecSet(Vec x, PetscScalar a);</code>	$x_i = a$

Some operations are *collective* while others are not.

PETSc Matrices

As with PETSc Vectors, there are many options for matrices:

- Those that are associated with individual processors
 - MatCreateSeqAIJ
 - MatCreateSeqSBAIJ
 - MatCreateSeqDense
 - MatCreateSeqBDiag

and

- those associated with an MPI communicator
 - MatCreateMPIAIJ
 - MatCreateMPISBAIJ
 - MatCreateMPIDense
 - MatCreateMPIBDiag

PETSc Matrices (2)

The preallocation of storage for sparse matrix data structures is essential for good performance.

Thus, we want a good estimate of the number of “diagonal” and “off-diagonal” nonzero entries for each row.

One useful tool is

```
MatSetOption(Mat A, MAT_NO_NEW_NONZERO_LOCATIONS);
```

when matrices keep the same nonzero pattern (in many time-dependent problems, nonlinear problems).

PETSc Matrices (3)

As with Vectors, `MatSetValues` is one tool for creating matrix entries.

- Again, either *insert* values

```
MatSetValues(Mat A, int m, int *rowidx, int n, int  
*colidx, PetscScalar *values, INSERT_VALUES);
```

- or *add* values

```
MatSetValues(Mat A, int m, int *rowidx, int n, int  
*colidx, PetscScalar *values, ADD_VALUES);
```

Once all components have been set with `MatSetValues()`, the following functions must be called

- `MatAssemblyBegin(Mat A, MAT_FINAL_ASSEMBLY);`
- `MatAssemblyEnd(Mat A, MAT_FINAL_ASSEMBLY);`

PETSc Matrices (4)

A partial list of functions on PETSc Matrices follows:

<code>MatAXPY(Mat Y, Mat X, a, MatStructure);</code>	$Y = Y + a * X$
<code>MatConvert(Mat A, MatType type, Mat B);</code>	$B = A$
<code>MatMult(Mat A, Vec x, Vec y);</code>	$y = A * x$
<code>MatNorm(Mat A, NormType type, double *r);</code>	$r = \ A\ _{type}$
<code>MatScale(Mat X, PetscScalar a);</code>	$X = a * X$
<code>MatTranspose(Mat A, Mat *B);</code>	$B = A^T$
<code>MatGetDiagonal(Mat A, Vec x);</code>	$x = \text{diag}(A)$
<code>MatZeroEntries(Mat A);</code>	$A = 0$
<code>MatShift(Mat X, PetscScalar a);</code>	$X = X + a * I$

PETSc Matrices (5)

Functions known as PETSc Viewers can be used to inspect matrices

- `MatView(Mat M, PETSC_VIEWER_STDOUT_WORLD);`

and

- `MatView(Mat M, PETSC_VIEWER_DRAW_WORLD);`

PETSc KSP: Linear System Solvers

Krylov subspace methods and preconditioners are the essential tool for solving *sparse* linear systems of the form

$$Ax = b.$$

This is setup with the commands

- `KSPCreate(PETSC_COMM_WORLD, KSP *ksp);`
- `KSPSetOperators(KSP ksp, Mat Amat, Mat Pmat, MatStructure flag);`

`Pmat` is the matrix used to create the preconditioner (usually `Amat`).

PETSc KSP: Linear System Solvers (2)

Many built-in Krylov subspace methods are available, including

Richardson	KSPRICHARDSON
Chebyshev	KSPCHEBYCHEV
Conjugate Gradient	KSPCG
BiConjugate Gradient	KSPBCG
Generalized Minimal Residual	KSPGMRES
BiCGSTAB	KSPBCGS
Conjugate Gradient Squared	KSPCGS
Transpose-Free Quasi-Minimal Residual	KSPTFQMR or KSPTCQMR

PETSc KSP: Linear System Solvers (3)

Either left or right preconditioning is possible (all of the built in preconditioners are left preconditioners by default).

$$(M_L^{-1} A M_R^{-1}) (M_R x) = (M_L^{-1} b)$$

The preconditioner is created using commands

- `KSPGetPC(KSP ksp, PC *pc);`
- `PCSetType(PC pc, PCtype option);`

where built in options include

Jacobi	PCJACOBI
Block Jacobi	PCBJACOBI
SOR	PCSOR
SOR with Eisenstat trick	PCEISENSTAT
Incomplete Cholesky	PCICC
Incomplete LU	PCILU

PETSc KSP: Linear System Solvers (4)

Specify Krylov subspace tolerances:

```
KSPSetTolerances(KSP ksp, double rtol, double atol,  
                 double dtol, int maxiter);
```

where

$$\|r_k\|_2 < \max(\text{rtol}\|b\|_2, \text{atol})$$

$k < \text{maxiter}$ corresponds to convergence, and

$$\|r_k\|_2 > \text{dtol}\|b\|_2$$

or $k = \text{maxiter}$ is divergence.

PETSC_DEFAULT can be used for any of these and can be overwritten at the command line with

```
KSPSetFromOptions(ksp);
```

First Example `ex1.c`

Solve

$$-u_{xx} = 0$$

on $(0, 1)$ with $u(0) = 1$ and $u(1) = 1$ using finite differences.

Then, partition $[0, 1]$ with points x_{-1}, \dots, x_n with

$$x_{i+1} - x_i = \Delta x = \frac{1}{n+1}$$

Thus, we set

$$-u_{xx}(x_i) \approx \frac{-u_{i-1} + 2u_i - u_{i+1}}{\Delta x} = 0$$

for $i = 0, \dots, n-1$ (knowing $u_{i-1} = 1 = u_n$).

First Example ex1.c

$$\begin{bmatrix}
 2 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\
 -1 & 2 & -1 & 0 & 0 & 0 & 0 & 0 \\
 0 & -1 & 2 & -1 & 0 & 0 & 0 & 0 \\
 & & \ddots & \ddots & \ddots & & & \\
 & & & \ddots & \ddots & \ddots & & \\
 & & & & \ddots & \ddots & \ddots & \\
 0 & 0 & 0 & 0 & 0 & -1 & 2 & -1 \\
 0 & 0 & 0 & 0 & 0 & 0 & -1 & 2
 \end{bmatrix}
 \begin{bmatrix}
 u_0 \\
 u_2 \\
 u_3 \\
 \vdots \\
 \vdots \\
 \vdots \\
 u_{n-2} \\
 u_{n-1}
 \end{bmatrix}
 =
 \begin{bmatrix}
 1 \\
 0 \\
 0 \\
 \vdots \\
 \vdots \\
 \vdots \\
 0 \\
 1
 \end{bmatrix}.$$

First Example ex1.c

```
#include "petscksp.h"

int main(int argc, char **args)
{
    Vec          u, b, exact;    // approx sol., RHS
    Mat          A;              // linear system matrix
    KSP          ksp;            // linear solver context
    PC           pc;             // preconditioner context
    PetscErrorCode ierr;
    PetscInt      i, n = 10, row[2], col[3], its;
    PetscReal     norm;
    PetscScalar   neg_one = -1.0, one = 1.0, value[3];

    PetscInitialize(&argc, &args, PETSC_NULL, PETSC_NULL);
    PetscOptionsGetInt(PETSC_NULL, "-n", &n, PETSC_NULL);
```

First Example ex1.c

```
VecCreateSeq(PETSC_COMM_SELF, n, &u);
VecDuplicate(u, &b);
VecDuplicate(u, &exact);

MatCreateSeq(PETSC_COMM_SELF, n, n, 0, PETSC_NULL, &A);

i = 0; col[0] = 0; col[1] = 1;
value[0] = 2.0; value[1] = -1.0;

MatSetValues(A,1,&i,2,col,value,INSERT_VALUES);
value[0] = -1.0; value[1] = 2.0; value[2] = -1.0;
for (i=1; i<n-1; i++) {
    col[0] = i-1; col[1] = i; col[2] = i+1;
    MatSetValues(A,1,&i,3,col,value,INSERT_VALUES);
}
```

First Example ex1.c

```
i = n - 1; col[0] = n - 2; col[1] = n - 1;  
value[0] = -1.0; value[1] = 2.0;  
MatSetValues(A,1,&i,2,col,value,INSERT_VALUES);
```

```
ierr = MatAssemblyBegin(A,MAT_FINAL_ASSEMBLY);CHKERRQ(ierr);  
ierr = MatAssemblyEnd(A,MAT_FINAL_ASSEMBLY);CHKERRQ(ierr);
```

```
row[0] = 0; value[0] = 1;  
row[1] = n-1; value[1] = 1;  
VecSetValues(b,2,row,value,INSERT_VALUES);
```

```
VecAssemblyBegin(b);  
VecAssemblyEnd(b);
```

```
VecSet(exact,one);
```

First Example `ex1.c`

```
KSPCreate(PETSC_COMM_WORLD,&ksp);
KSPSetOperators(ksp,A,A,DIFFERENT_NONZERO_PATTERN);
KSPGetPC(ksp,&pc);
PCSetType(pc,PCJACOBI);
KSPSetTolerances(ksp,1.e-7,PETSC_DEFAULT,
                 PETSC_DEFAULT,PETSC_DEFAULT);
KSPSolve(ksp,b,u);

VecAXPY(u,neg_one,exact);
VecNorm(u,NORM_2,&norm);

KSPGetIterationNumber(ksp,&its);
PetscPrintf(PETSC_COMM_WORLD,
            "Norm of error %A, Iterations %D\n",
            norm,its);
```

First Example ex1.c

```
VecDestroy(exact);  
VecDestroy(b);  
VecDestroy(u);  
  
MatDestroy(A);  
  
KSPDestroy(ksp);  
  
PetscFinalize();  
return 0;  
}
```

Typing `mpirun -np 1 ex1` leads to the output:

Norm of error < 1.e-12, Iterations 5

Second Example ex23.c

```
#include "petscksp.h"

int main(int argc,char **args)
{
    Vec          u, b, exact;    // approx sol., RHS, exact
    Mat          A;              // linear system matrix
    KSP           ksp;           // linear solver context
    PC            pc;            // preconditioner context
    PetscInt      i, n = 10000, row[1], col[3], its,
                 rstart, rend, nlocal;

    PetscReal     norm;          // norm of solution error
    PetscScalar   neg_one = -1.0, one = 1.0, value[3];

    PetscInitialize(&argc,&args,PETSC_NULL,PETSC_NULL);
    PetscOptionsGetInt(PETSC_NULL,"-n",&n,PETSC_NULL);
```

Second Example ex23.c

```
PreLoadBegin(PETSC_TRUE,"beginning of code");
```

```
VecCreate(PETSC_COMM_WORLD,&u);
```

```
VecSetSizes(u,PETSC_DECIDE,n);
```

```
VecSetFromOptions(u);
```

```
VecDuplicate(u,&b);
```

```
VecDuplicate(u,&exact);
```

```
VecGetOwnershipRange(u,&rstart,&rend);
```

```
VecGetLocalSize(u,&nlocal);
```

```
MatCreate(PETSC_COMM_WORLD,&A);
```

```
MatSetSizes(A,nlocal,nlocal,n,n);
```

```
MatSetFromOptions(A);
```

Second Example ex23.c

```
if (rstart == 0) {
    rstart = 1;
    i = 0; col[0] = 0; col[1] = 1; value[0] = 2.0; value[1]
    MatSetValues(A,1,&i,2,col,value,INSERT_VALUES);

    row[0] = 0; value[0] = 1.0;
    VecSetValues(b,1,row,value,INSERT_VALUES);
}

if (rend == n) {
    rend = n-1;
    i = n-1; col[0] = n-2; col[1] = n-1; value[0] = -1.0; va
    MatSetValues(A,1,&i,2,col,value,INSERT_VALUES);

    row[0] = n-1; value[0] = 1.0;
    VecSetValues(b,1,row,value,INSERT_VALUES);
}
```

Second Example ex23.c

```
value[0] = -1.0; value[1] = 2.0; value[2] = -1.0;
for (i=rstart; i<rend; i++) {
    col[0] = i-1; col[1] = i; col[2] = i+1;
    MatSetValues(A,1,&i,3,col,value,INSERT_VALUES);
}
```

```
MatAssemblyBegin(A,MAT_FINAL_ASSEMBLY);
MatAssemblyEnd(A,MAT_FINAL_ASSEMBLY);
```

```
VecAssemblyBegin(b);
VecAssemblyEnd(b);
```

```
VecSet(exact,one);
```

Second Example ex23.c

```
KSPCreate(PETSC_COMM_WORLD,&ksp);
KSPSetOperators(ksp,A,A,DIFFERENT_NONZERO_PATTERN);

KSPGetPC(ksp,&pc);
PCSetType(pc,PCJACOBI);
KSPSetTolerances(ksp,1.e-7,PETSC_DEFAULT,PETSC_DEFAULT,PETSC_DEFAULT);
KSPSetFromOptions(ksp);
KSPSolve(ksp,b,u);
KSPView(ksp,PETSC_VIEWER_STDOUT_WORLD);
VecAXPY(u,neg_one,exact);
VecNorm(u,NORM_2,&norm);
KSPGetIterationNumber(ksp,&its);

PetscPrintf(PETSC_COMM_WORLD,
            "Norm of error %A, Iterations %D\n",
            norm,its);
```

Second Example ex23.c

```
VecDestroy(exact);  
VecDestroy(b);  
VecDestroy(u);  
MatDestroy(A);  
KSPDestroy(ksp);  
  
PreLoadEnd();  
PetscFinalize();  
return 0;  
}
```

Typing `mpirun -np 1 ex23` or `mpirun -np 4 ex23` each lead to the same output:

Norm of error < 1.e-12, Iterations 5

ViTLES - Virginia Tech Large-Eddy Simulator

Developed with Traian Iliescu, uses PETSc to solve 2D/3D LES model equations.

Runs on System X: www.tcf.vt.edu



ViTLES - Virginia Tech Large-Eddy Simulator


ViTLES

The Virginia Tech Large-Eddy Simulator

- Subgrid scale models
 - dynamic
 - deconvolution
 - differential filters
- Boundary condition models
 - Approximate deconvolution boundary conditions
- Reduced-order modeling
- Control

ViTLES - Virginia Tech Large-Eddy Simulator

Finite elements: requires different PETSc functions



```
VecScatterCreate(g_ctx->crt_N_iter,
                 l_ctx->from_is,
                 l_ctx->g_vec,
                 l_ctx->to_is, &scatter);
VecScatterBegin(g_ctx->crt_N_iter,
                 l_ctx->g_vec,
                 INSERT_VALUES,
                 SCATTER_FORWARD,
                 scatter);
VecScatterEnd(g_ctx->crt_N_iter,
               l_ctx->g_vec,
               INSERT_VALUES,
               SCATTER_FORWARD,
               scatter);
```